

File management and manipulation

Операционни системи, ФМИ, 2017/2018

Filesystem Hierarchy Standard

- Filesystem standard (FHS)
 - Guiding principles for each area of filesystem
 - Predictable location of files and directories
 - Provides uniformity across multiple Linux distributions
- The Linux Standards Base
 - Aims to allow Linux binaries to run unmodified on multiple Linux distributions
 - Specifies system and library interfaces and environment
 - Incorporates the FHS

Navigating the filesystem

- Absolute vs. relative addressing
- Changing and displaying directories (cd, pwd)
- cd (without parameters)
- cd ~george
- cd ~
- cd -
- . and ..

Displaying directory contents

- human-readable
- `ls`
- `ls -a` - show all files (including *.hidden* files)
- `ls -l` - long listings
- `ls -d` - show directories, not contents
- `touch foo, mkdir bar`


File {group,}ownership

- each file is owned by a specific UID and GID
- `chown` - change the user (UID) ownership
 - only root can change ownership to another user
 - `chown foo:bar`
- `chgrp` - modify just the group (GID) ownership
- newly created files will usually be given GID ownership based on the current active group of the person who creates the file
- `newgrp foo` - log in to a new group
 - newly created files will be owned by the new group
 - users can only change to their own groups
 - root user can change to any group
 - `exit` to switch back

File permissions

- type of file
 - - - regular file
 - b - block special file
 - c - character special file
 - d - directory
 - l - symbolic link
 - p - FIFO (named pipe)
 - s - socket
- permission sets
 - user (owner)
 - group (group owner)
 - everyone else (other)
 - symbolic representation `rw-r--r--`
 - numeric representation 0755

File permissions (cont.)



- r - 100b - 4 - Read
- w - 010b - 2 - Write
- x - 001b - 1 - Execute

Special permissions

- Set UID upon execution (SUID)
- Set GID upon execution (SGID)
- sticky bit
- different behavior for files and directories




Figure 1: SUID

Special permissions (cont.)

- SUID and SGID on files
 - An executable with the SUID bit set runs with the security context of the user who owns it, regardless of the executing user
 - SGID
- SGID on directories
 - Files or sub-directories created within that directory inherit the group ownership of the SGID directory
- Sticky Bit on directories
 - Normally in a directory that is world writable, users can delete each other's files. Setting the sticky bit overrides this behavior

Changing file permissions

- `chmod`
 - numeric notation `chmod 0664 foo.txt`
 - symbolic notation `chmod u=rw,g=rw,o=r foo.txt`
 - `+`, `-`, `=`
- `chmod -R`

umask

- Default permissions for newly created filesystem objects
 - files 666
 - directories 777
- umask
 - defines what permissions to **withhold** from the default permissions
 - display or change umask
 - usually set in the user or system shell dot files

User Private Group (UPG) scheme

- convenient way to share files when working in a group project directory
- each user in their own private group
- `umask 0002`
- set GID of project directory to a commonly shared GID
- set SGID on the project directory


Directory and file manipulation

- `mkdir foo`
 - `mkdir -p foo/bar`
 - `mkdir -m`
- `rmdir`
- `cp`
- `mv`
- `rm`
- `touch -mtime/ctime`

UNIX filesystem structure

- blocks
- inodes
 - permissions
 - access time, modification time, inode change time
 - owner
 - group
 - size in bytes
 - occupied blocks
 - link count (names of the inode)
 - inode number
- directories (are files that) hold filenames and inodes
- superblock contains filesystem parameters (how many inodes, etc.)

inode pointer structure



Filesystem hard links

- a directory entry that references the same inode as another directory entry
 - can't span filesystems
 - can't create hard links to non-existent file
 - can't reference directories
 - do not occupy storage space (i.e. blocks)
 - `ln [option]... target link_name`
- `ls -li`





Figure 3: hardlinks

Filesystem symbolic links

- a file that references another file via path and name
 - can reference directories
 - can span filesystems
 - can reference non-existent files
 - `ln -s target link_name`
 - occupy space
- symlink / soft link



df, du, stat

- df Report disk space usage per filesystem
 - -h human readable output
 - -i list inode information instead of block usage
 - -T include filesystem type
 - --si use powers of 1000 instead of 1024
 - -P use the POSIX output format
- du Report disk usage per file and directory
 - -h human readable sizes
 - -s summarize, only display total for each argument
 - -x do not include files on a different filesystem
 - --si use powers of 1000 instead of 1024
- stat display file or file system status
 - -L follow links
 - -c --format

File extensions and content

- file extensions are just part of the file name
- some applications may care about extensions
- `file` - reports the type of file by examining the file contents
- `/usr/share/file/magic.mgc`

Displaying text files

- `cat` - concatenate files and print on the standard output
- `more`
- `less`
- `head`
- `tail`
 - `tail -f`
- `-n`

Displaying binary files

- Displaying raw binary data may corrupt the display terminal
- `reset`
- `[Ctrl-J] reset [Ctrl-J]` - if carriage-return fails
- `strings` - displays ASCII text inside binary files
- `xxd` - displays HEX and ASCII dump of file
- `clear`

Searching the filesystem

- machine-readable
- `find [options] [starting-point] [expression]`
- starting point
- criteria
 - `-print` vs `-print0` vs `-printf`
 - `-ls`
 - `-exec`
- `find /foo -name bar -print`

Archiving & compressing

- archiving
 - tar
 - cpio
- compressing
 - compress
 - gzip
 - bzip2
 - lzma
 - xz

Archives with tar

- tar
 - manipulates `.tar` files (*tarballs*)
 - used for backup and transfer of files
 - creates, extracts or lists the contents of tarballs
 - `c`, `x`, `t`, `f`, `v`
 - traditional vs. UNIX-style vs. GNU-style usage
 - `tar cvf foo.tar ./foo/*`
 - GNU tar supports built-in compression methods
- `.tar` (*tarball*)
 - records file and directory structure
 - includes metadata about the file: date, timestamps, ownership, permissions, etc.

Archives with cpio

- manipulates .cpio files
- used as the basis for RPM packages
- doesn't recurse sub-directories, must be passed list of directories
- more robust than tar when media errors encountered
- -i input mode, used when feeding a cpio archive into the cpio command
- -o output mode, used to create cpio archives, which are sent to STDOUT

The gzip compression utility

- `gzip` popular replacement for `compress`
 - created by the GNU project because of patented algorithms in `compress`
 - default action deletes original after creating new compressed file
 - standard file extension: `.gz`
 - much higher compression ratio than `compress`
- `gunzip` or `zcat` decompress `.gz` files
 - `gunzip` decompresses the file on disk (removing the original, compressed file); `zcat` does not
 - `zcat` outputs uncompressed data to `STDOUT`

The bzip2 compression utility

- bzip2 typically better compression than gzip
- .bz2
- bunzip2 / bzip2
- replaces gzip as compression format of choice

XZ Utils

- xz typically better compression than bzip2
- .xz
- unxz / xzcat
- replaces bzip2 as compression format of choice